

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Full-Stack Founding Engineer with end-to-end ownership across React/React Native frontends, Python (FastAPI/Django) and Node.js backends, PostgreSQL/Redis data layers, and real-time WebSocket systems. Specialised in shipping production LLM agents (OpenAI, Gemini, Groq) with provider-fallback, prompt engineering, and RAG pipelines. Proven track record: 40% API latency reduction, 30% faster deployment cycles, 40+ SMB pilots, 1000+ users served.

PROFESSIONAL EXPERIENCE

OpenBiz Software India Pvt Ltd

Jun 2025 – May 2026

Founding Engineer

Remote/Bangalore

- Architected the zero-to-one core MERN platform with strict TypeScript practices, building a scalable, production-ready foundation serving 1000+ users.
- Built VyaparGPT: WhatsApp-native LLM assistant for Indian SMBs – webhook ingestion → LLM orchestration → response delivery, deployed to 40+ closed-pilot businesses.
- LLM provider-fallback: OpenAI → Google Gemini on HTTP 5xx/429/timeout with response-shape normalisation – zero failure surface exposed to users.
- Engineered real-time bidirectional chat infrastructure using Socket.IO enabling live messaging between users.
- Diagnosed and resolved N+1 query patterns in Supabase (PostgreSQL), reducing API latency by 40% through indexing and response caching.
- Owned CI/CD pipeline via GitHub Actions and Vercel, reducing deployment cycle time by 30% and eliminating environment-drift bugs.
- Engineered fault-tolerant Razorpay subscription billing with idempotent webhook validation, achieving 100% payment-state consistency.

Vellore Institute of Technology

Jan 2025 – Nov 2025

Research Contributor

Vellore

- Contributed to research on Blockchain-Based LLM Model Using Fully Homomorphic Encryption (FHE) for Academic Records.

PROJECTS

Nimbus – AI RAG Workspace | github.com/atinsharma24/nimbus

In Development 2026

- Stack: Next.js, pgvector, Groq, Prisma
- Architected a production RAG pipeline (pgvector + Groq) achieving sub-800ms semantic retrieval over multi-format corpora with streaming LLM responses.
- Provider-agnostic LLM layer supporting Groq, OpenAI, and Anthropic; async ingestion with 429 retry logic and per-document status tracking.

Autonomous Web Agent | github.com/atinsharma24/auto-agent

- Stack: OpenClaw, BrowserOS (MCP), Python, Groq
- Claude 3.5 Sonnet + MCP orchestration layer translating structured context schemas into DOM-level form interactions for multi-step web workflows.
- Decoupled Groq-based “Watcher” (cost-efficient parsing) + heavy “Executor” agent; local queue management with WAF/Captcha detection.

DocGPT – RAG Document Assistant | github.com/atinsharma24/docgpt

- Stack: React 19, Python (Django), ChromaDB, OpenAI, Docker
- Full Django RAG pipeline: PyPDF2 chunking, local all-MiniLM-L6-v2 embeddings (zero API cost), ChromaDB cosine retrieval, GPT-4o-mini generation.
- Containerized Django + ChromaDB as Docker services for reproducible local deployment.

SKILLS

Focus Stack: React 19, React Native, TypeScript, Node.js, Python/FastAPI, PostgreSQL, Redis, WebSockets, LLM Agents, Docker

Languages: JavaScript/TypeScript, Python, SQL, Java

Frontend: React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

Backend/DB: Node.js, Express, NestJS, FastAPI, Django, PostgreSQL, Redis, MongoDB, pgvector, ChromaDB, WebSockets

AI/LLM: OpenAI, Google Gemini, Groq API, RAG Pipelines, LangChain.js, Sentence Transformers, Prompt Engineering, LLM Agents

DevOps: Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB)

Tools: Git, Postman, Socket.IO, Razorpay SDK, VideoSDK, Puppeteer

EDUCATION

Vellore Institute of Technology

Expected Graduation: May 2026

B.Tech in Computer Science and Engineering

Vellore

- Focus on Full Stack Development and AI Systems.