

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Specialized Generative AI Engineer focused on LLM orchestration, agentic workflows, and high-fidelity RAG pipelines. Expert in prompt engineering, dual-model failover architectures (Gemini + OpenAI), and semantic search optimization using pgvector and ChromaDB. Experienced in designing resilient AI infrastructures with sub-800ms retrieval times and fault-tolerant message lifecycles.

- Architected Nimbus AI: production RAG workspace with provider-agnostic LLM layer + async ingestion.
- Built Autonomous Web Agent: Claude 3.5 Sonnet + MCP orchestration for DOM-level form interaction.
- Content Intelligence Pipeline: Gemini → OpenAI dual-model rewriting + React 19 Diff View.

PROFESSIONAL EXPERIENCE

OpenBiz Software India Pvt Ltd

Jun 2025 – May 2026

Founding Engineer

Remote / Bangalore

- VyaparGPT: HMAC-SHA256 webhook → session context → LLM dispatch → delivery; 40+ SMBs pilot, 1,000+ platform users; multi-turn context via `conversation_history` [].
- LLM fallback (OpenAI → Gemini on 5xx/429/timeout); Gemini Vision extracted GST, amounts, dates from WhatsApp SMB photographs, eliminating manual onboarding data entry.
- N+1 fix (JOIN rewrites, composite indexes, TTL cache) – **40% API latency reduction** at 1,000+ user scale.
- Razorpay: HMAC webhook validation + re-entrant idempotency check – **100% payment-state consistency**.
- CI/CD (GitHub Actions + Vercel): automated testing per PR, staging on merge, prod behind approval – **30% faster deployments**, zero broken-build incidents over four months.

PROJECTS

Nimbus – AI RAG Workspace | Next.js 14, TypeScript, PostgreSQL, pgvector, Prisma, Groq API, OpenAI

In Development 2026

- Async queue ingestion (pending → indexed) with overlap chunking; pgvector cosine search for hybrid SQL + semantic filtering – sub-800ms retrieval.
- Provider-agnostic LLM layer (groq|openai|anthropic config flag) + SSE streaming via Groq; 429 retry (exp. backoff) and per-document status tracking.

DocGPT – RAG Document Assistant | React 19, Python (Django), ChromaDB, OpenAI, Docker

- Django RAG pipeline: PyPDF2 overlap chunking, local all-MiniLM-L6-v2 embeddings (offline, zero API cost), ChromaDB cosine retrieval, GPT-4o-mini generation; Dockerized.

Autonomous Web Agent | github.com/atinsharma24/auto-agent | Python, Groq, BrowserOS (MCP), Claude 3.5 Sonnet

- Dual-agent: Claude 3.5 Sonnet Executor (DOM + form reasoning) + Groq Watcher (page-state validation between transitions).
- MCP via BrowserOS for structured DOM access; persistent queue (pending → applied/failed) with URL dedup and WAF/CAPTCHA detection.

Automated Content Intelligence Pipeline | React 19, TypeScript, Laravel, OpenAI, Gemini, Docker, Puppeteer

- Puppeteer SPA scraping; Gemini → OpenAI fallback rewriting; React 19 Diff View with per-article state machine; Dockerized microservices on GitHub Actions cron.

SKILLS

Focus Stack: Python, LangChain, Groq API, pgvector, ChromaDB, OpenAI, Google Gemini, Prompt Engineering, Agentic Workflows.

Languages: JavaScript/TypeScript, Python, SQL, Java

Frontend: React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

Backend/DB: Node.js, Express, NestJS, Django, PostgreSQL, MongoDB, pgvector, ChromaDB

AI/LLM: OpenAI, Google Gemini, Groq API, RAG Pipelines, LangChain.js, Sentence Transformers, Prompt Engineering

DevOps: Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB)

Tools: Git, Postman, Socket.IO, Razorpay SDK, VideoSDK, Puppeteer

Patterns: Webhook Validation (HMAC-SHA256), N+1 Optimization, Vector Similarity Search, Provider-Agnostic LLM Layers, Idempotent Events

EDUCATION

Vellore Institute of Technology

2026

B.Tech in Computer Science and Engineering

Vellore, India

- Research (2025): Blockchain-based credential audit trail with FHE-enabled encrypted-data LLM inference.