

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Full-Stack Engineer with end-to-end production ownership across Next.js/React frontends, Node.js and Python (FastAPI/Django) backends, PostgreSQL/Redis data layers, and real-time WebSocket systems. Specialised in production LLM agents (OpenAI, Gemini, Groq, Claude) with provider failover, prompt engineering, and RAG pipelines, built through an AI-native workflow with Claude Code and Codex. Track record: 40% API latency reduction, 30% faster deployments, 40+ SMB pilots, 1,000+ users served.

PROFESSIONAL EXPERIENCE

AGI Ready Inc.

Jun 2026 – Present

Full-Stack Developer

Remote / Delaware, US

- Build and ship full-stack product features in Next.js, TypeScript, and Node.js across multiple concurrent client products, working directly with the founder and a small delivery team.
- Architect and integrate service APIs across third-party platforms, AI/LLM providers, webhook receivers, and asynchronous background workflows.
- Operate an AI-native development workflow with Claude Code and Codex – directing agentic tooling with architectural judgement rather than prompt-level iteration.
- Own production systems end-to-end – debugging, deployment, and release – decomposing ambiguous client problems into scoped deliverables shipped iteratively.

OpenBiz Software India Pvt Ltd

Jun 2025 – May 2026

Founding Engineer

Remote / Bangalore

- Architected the zero-to-one MERN platform in strict TypeScript, scaling to 1,000+ active users; engineered a parallel Graph RAG pipeline with LangGraph state machines routing queries into async multi-threaded retrieval paths.
- Built a dual-database ingestion router (Drizzle ORM) streaming large-scale datasets into Pinecone (hybrid dense/sparse) and Supabase/pgvector – partitioning unstructured from structured compliance data.
- Cut API latency 40% at 1,000+ user scale by resolving N+1 chains with JOIN rewrites, composite indexes, and TTL caching; built fault-tolerant Razorpay billing with HMAC-SHA256 idempotent webhooks for 100% payment-state consistency.
- Owned the CI/CD pipeline (GitHub Actions + Vercel) with per-PR automated testing, staged merges, and approval-gated production – cutting deployment time by 30% with zero broken-build incidents.

Vellore Institute of Technology

Jan 2025 – Nov 2025

Research Contributor

Vellore

- Research on blockchain-based LLM integration with Fully Homomorphic Encryption (FHE) for encrypted-data inference without decryption.

PROJECTS

Nimbus – AI RAG Workspace | github.com/atinsharma24/nimbus

In Development

- Stack: Next.js 14, TypeScript, PostgreSQL, pgvector, Prisma, Groq API, OpenAI
- Engineered a production RAG pipeline (pgvector + Groq) delivering sub-800ms semantic retrieval over multi-format corpora with SSE streaming.
- Provider-agnostic LLM layer (Groq/OpenAI/Anthropic) with async queue ingestion, exponential-backoff 429 retry, and per-document status tracking.

Autonomous Web Agent | github.com/atinsharma24/auto-agent

- Stack: Python, Groq, BrowserOS (MCP), Claude 3.5 Sonnet
- Orchestrated a dual-agent architecture: a Claude 3.5 Sonnet Executor for DOM/form reasoning plus a Groq Watcher for page-state validation between steps.
- Integrated BrowserOS MCP for structured DOM access; persistent job queue with URL deduplication and WAF/CAPTCHA detection.

DocGPT – RAG Document Assistant | github.com/atinsharma24/docgpt

- Stack: React 19, Python (Django), ChromaDB, OpenAI, Docker
- Built an end-to-end Django RAG pipeline: PyPDF2 overlap chunking, local all-MiniLM-L6-v2 embeddings (zero API cost), ChromaDB cosine retrieval, and GPT-4o-mini generation.
- Containerised Django + ChromaDB as independent Docker services for reproducible local deployment.

SKILLS

Focus Stack: Next.js, React 19, TypeScript, Node.js, Python/FastAPI, PostgreSQL, Redis, LLM Agents, Docker

Languages & Frontend: JavaScript/TypeScript, Python, SQL, Java | React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

Backend & Data: Node.js, Express, NestJS, FastAPI, Django, PostgreSQL, Supabase, Redis, MongoDB, pgvector, ChromaDB, Pinecone, WebSockets

AI/LLM: OpenAI, Google Gemini, Groq, Anthropic Claude, RAG Pipelines, LangChain.js/LangGraph, Sentence Transformers, Prompt Engineering, MCP, Claude Code, Codex

DevOps & Tools: Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB), Git, Postman, Socket.IO, Razorpay SDK, Puppeteer

EDUCATION

Vellore Institute of Technology

B.Tech in Computer Science and Engineering – Focus: Full-Stack Development & AI Systems

2026

Vellore, India