

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Specialized Generative AI Engineer focused on LLM orchestration, agentic workflows, and high-fidelity RAG pipelines. Expert in prompt engineering, dual-model failover architectures (Gemini + OpenAI), and semantic search optimization using pgvector and ChromaDB. Experienced in designing resilient AI infrastructures with sub-800ms retrieval times and fault-tolerant message lifecycles.

PROFESSIONAL EXPERIENCE

AGI Ready Inc.

Jun 2026 – Present

Full-Stack Developer

Remote / Delaware, US

- Operate an AI-native development workflow with Claude Code and Codex – decomposing problems, directing agentic tooling with architectural judgement, and reviewing generated output against production standards.
- Architect and integrate AI/LLM provider APIs, webhook receivers, and asynchronous background workflows across multiple concurrent client products.
- Build and ship full-stack features in Next.js, TypeScript, and Node.js; own production systems end-to-end through debugging, deployment, and release.

OpenBiz Software India Pvt Ltd

Jun 2025 – May 2026

Founding Engineer

Remote / Bangalore

- Engineered a parallel Graph RAG pipeline using LangGraph state machines routing queries into async multi-threaded retrieval paths; built a dual-database ingestion router (Drizzle ORM) partitioning unstructured data into Pinecone (hybrid dense/sparse) and structured compliance data into Supabase (PostgreSQL/pgvector).
- Designed runtime evaluation guardrails intercepting Cohere Re-rank v3 cross-encoder scores; an absolute < 0.4 confidence threshold hook dynamically overrides system prompts to enforce strict context-locking and suppress hallucination.
- Built VyaparGPT: WhatsApp-native LLM assistant with automated failover (OpenAI → Gemini) and Gemini Vision OCR for invoice processing.
- CI/CD (GitHub Actions + Vercel): automated testing per PR, staging on merge, prod behind approval – **30% faster deployments**, zero broken-build incidents over four months.

PROJECTS

Nimbus – AI RAG Workspace | github.com/atinsharma24/nimbus

In Development 2026

- Stack: Next.js 14, TypeScript, PostgreSQL, pgvector, Prisma, Groq API, OpenAI
- Async queue ingestion (pending → indexed) with overlap chunking; pgvector cosine search for hybrid SQL + semantic filtering – sub-800ms retrieval.
- Provider-agnostic LLM layer (groq|openai|anthropic config flag) + SSE streaming via Groq; 429 retry (exp. backoff) and per-document status tracking.

DocGPT – RAG Document Assistant | github.com/atinsharma24/docgpt

- Stack: React 19, Python (Django), ChromaDB, OpenAI, Docker
- Django RAG pipeline: PyPDF2 overlap chunking, local all-MiniLM-L6-v2 embeddings (offline, zero API cost), ChromaDB cosine retrieval, GPT-4o-mini generation; Dockerized.

Autonomous Web Agent | github.com/atinsharma24/auto-agent

- Stack: Python, Groq, BrowserOS (MCP), Claude 3.5 Sonnet
- Dual-agent: Claude 3.5 Sonnet Executor (DOM + form reasoning) + Groq Watcher (page-state validation between transitions).
- MCP via BrowserOS for structured DOM access; persistent queue (pending → applied/failed) with URL dedup and WAF/CAPTCHA detection.

Automated Content Intelligence Pipeline | github.com/atinsharma24/pipeline

- Stack: React 19, TypeScript, Laravel, OpenAI, Gemini, Docker, Puppeteer
- Puppeteer SPA scraping; Gemini → OpenAI fallback rewriting; React 19 Diff View with per-article state machine; Dockerized microservices on GitHub Actions cron.

SKILLS

Focus Stack: Python, LangChain/LangGraph, Groq API, pgvector, ChromaDB, OpenAI, Google Gemini, Anthropic Claude, Prompt Engineering, Agentic Workflows

AI/LLM: OpenAI, Google Gemini, Groq, Anthropic Claude, RAG Pipelines, LangChain.js, Cohere Re-rank, Sentence Transformers, MCP, Claude Code, Codex

Languages & Frontend: JavaScript/TypeScript, Python, SQL, Java | React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

Backend & Data: Node.js, Express, NestJS, Django, FastAPI, PostgreSQL, Supabase, MongoDB, pgvector, ChromaDB, Pinecone

DevOps, Tools & Patterns: Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB), Git, Puppeteer | HMAC-SHA256 webhook validation, provider-agnostic LLM layers, vector similarity search, idempotent events

EDUCATION

Vellore Institute of Technology

B.Tech in Computer Science and Engineering – Research (2025): FHE-enabled encrypted-data LLM inference

2026

Vellore, India