

# ATIN SHARMA

Agra, India | [atinsharma24@gmail.com](mailto:atinsharma24@gmail.com) | +91 82185 02886  
[linkedin.com/in/atinsharma24](https://linkedin.com/in/atinsharma24) | [github.com/atinsharma24](https://github.com/atinsharma24)

## PROFILE

Full-Stack Founding Engineer with end-to-end ownership across React/React Native frontends, Python (FastAPI/Django) and Node.js backends, PostgreSQL/Redis data layers, and real-time WebSocket systems. Specialized in shipping production LLM agents (OpenAI, Gemini, Groq) with provider fallover, prompt engineering, and RAG pipelines. Proven track record: 40% API latency reduction, 30% faster deployment cycles, 40+ SMB pilots, 1,000+ users served.

## PROFESSIONAL EXPERIENCE

### OpenBiz Software India Pvt Ltd

Jun 2025 – May 2026

Founding Engineer

Remote / Bangalore

- Architected the zero-to-one MERN platform from scratch with strict TypeScript practices, delivering a production-ready foundation that scaled to 1,000+ active users.
- Engineered a parallel Graph RAG pipeline with LangGraph state machines, routing queries into async multi-threaded retrieval paths to reduce end-to-end token latency.
- Built a dual-database ingestion router (Drizzle ORM) streaming large-scale datasets into Pinecone (Hybrid Dense/Sparse) and Supabase/pgvector — partitioning unstructured from structured compliance data.
- Engineered fault-tolerant Razorpay subscription billing with HMAC-SHA256 idempotent webhook validation, achieving 100% payment-state consistency.
- Resolved N+1 query chains with JOIN rewrites, composite indexes, and TTL caching — cutting API latency by 40% at 1,000+ user scale.
- Owned CI/CD pipeline (GitHub Actions + Vercel) with per-PR automated testing, staged merges, and approval-gated production — cutting deployment time by 30% with zero broken-build incidents.

### Vellore Institute of Technology

Jan 2025 – Nov 2025

Research Contributor

Vellore

- Contributed to research on Blockchain-based LLM integration with Fully Homomorphic Encryption (FHE) for encrypted-data inference without decryption.

## PROJECTS

### Nimbus — AI RAG Workspace | [github.com/atinsharma24/nimbus](https://github.com/atinsharma24/nimbus)

In Development

- Stack: Next.js 14, TypeScript, PostgreSQL, pgvector, Prisma, Groq API, OpenAI
- Engineered production RAG pipeline (pgvector + Groq) delivering sub-800ms semantic retrieval over multi-format corpora with SSE streaming.
- Provider-agnostic LLM layer (Groq/OpenAI/Anthropic) with async queue ingestion, exponential-backoff 429 retry, and per-document status tracking.

### Autonomous Web Agent | [github.com/atinsharma24/auto-agent](https://github.com/atinsharma24/auto-agent)

- Stack: Python, Groq, BrowserOS (MCP), Claude 3.5 Sonnet
- Orchestrated dual-agent architecture: Claude 3.5 Sonnet Executor for DOM and form reasoning plus Groq Watcher for page-state validation between steps.
- Integrated BrowserOS MCP for structured DOM access; persistent job queue with URL deduplication and WAF/CAPTCHA detection.

### DocGPT — RAG Document Assistant | [github.com/atinsharma24/docgpt](https://github.com/atinsharma24/docgpt)

- Stack: React 19, Python (Django), ChromaDB, OpenAI, Docker
- Built end-to-end Django RAG pipeline: PyPDF2 overlap chunking, local all-MiniLM-L6-v2 embeddings (zero API cost), ChromaDB cosine retrieval, and GPT-4o-mini generation.
- Containerized Django + ChromaDB as independent Docker services for reproducible local deployment.

## SKILLS

**Focus Stack:** React 19, React Native, TypeScript, Node.js, Python/FastAPI, PostgreSQL, Redis, WebSockets, LLM Agents, Docker

**Languages:** JavaScript/TypeScript, Python, SQL, Java

**Frontend:** React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

**Backend/DB:** Node.js, Express, NestJS, FastAPI, Django, PostgreSQL, Redis, MongoDB, pgvector, ChromaDB, WebSockets

**AI/LLM:** OpenAI, Google Gemini, Groq API, RAG Pipelines, LangChain.js, Sentence Transformers, Prompt Engineering, LLM Agents

**DevOps:** Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB)

**Tools:** Git, Postman, Socket.IO, Razorpay SDK, VideoSDK, Puppeteer

## EDUCATION

### Vellore Institute of Technology

B.Tech in Computer Science and Engineering

Vellore, India

- Focus on Full-Stack Development and AI Systems.