

ATIN SHARMA

Full-Stack Product Engineer | AI Systems | RAG Pipelines | WhatsApp Automation
atinsharma24@gmail.com | +91 82185 02886 | github.com/atinsharma24/ | linkedin.com/in/atinsharma24/
Agra, Uttar Pradesh, India — Open to Remote & Relocation

PROFESSIONAL SUMMARY

Founding engineer with 12 months of production experience building AI-powered full-stack systems at OpenBiz Software India Pvt Ltd — one of three engineers who shipped two commercial products from zero to 1,000+ active Indian SMB users. Specializes in MERN + AI stacks (TypeScript, Node.js, Python, PostgreSQL) with hands-on production delivery across conversational AI (VyaparGPT, WhatsApp-native LLM assistant for 40+ SMB pilot businesses), production RAG pipelines (Nimbus on pgvector + Groq; DocGPT on ChromaDB), and resilient DevOps (GitHub Actions + Vercel CI/CD reducing deployment cycle time by 30%). Technical scope spans full-stack product engineering, dual-model LLM provider-fallback (OpenAI ↔ Gemini), idempotent Razorpay billing, real-time Socket.IO infrastructure, and autonomous multi-agent web automation (Claude 3.5 Sonnet + MCP). Immediate joiner (0-day notice) targeting 15–18 LPA in founding-stage or early-product roles with full-stack ownership. VIT Vellore, B.Tech Computer Science Engineering, graduating May 2026.

KEY METRICS

API Latency Reduction: 40%	Platform Active Users: 1,000+
Deployment Cycle Faster: 30%	VyaparGPT SMB Pilot: 40 businesses
Shared Infra Reuse (Website Builder): ~60%	LeetCode Problems Solved: 347
Founding Team Size: 3 engineers	Zero Broken-Build Incidents: 4 months post CI/CD

PROFESSIONAL EXPERIENCE

OpenBiz Software India Pvt Ltd | Founding Engineer | Jun 2025 – May 2026 | Remote / Bangalore

One of three founding technical leads; architected and shipped two AI-powered SaaS products for the Indian SMB market on the MERN stack.

- Architected **VyaparGPT**, a WhatsApp-native LLM assistant in Node.js: designed the full message lifecycle (WhatsApp Business API webhook ingestion → HMAC-SHA256 validation → session context lookup → LLM dispatch → response delivery), serving **40+ SMBs** in closed pilot and scaling to **1,000+ active platform users**. Multi-turn conversational context via rolling `conversation_history[]` enables natural follow-up queries.
- Engineered **LLM provider-fallback** (OpenAI → Google Gemini): on HTTP 5xx, 429 rate-limit, or configurable timeout, requests are transparently rerouted with response-shape normalization — zero failure surface exposed to users.
- Built a **Gemini Vision API document verification module** extracting structured fields (GST number, business name, invoice amount, dates) from WhatsApp photographs of GST certificates and vendor agreements, eliminating a manual data-entry step in SMB onboarding.
- Resolved critical **N+1 query patterns** in Supabase (PostgreSQL): rewrote ORM calls to JOIN-based queries, introduced composite indexes on high-frequency WHERE/JOIN columns, and added TTL-based response caching — achieving verified **40% API latency reduction** at 1,000+ user scale.
- Built and owned **GitHub Actions + Vercel CI/CD pipeline** from scratch: automated testing on every PR, staging deploys on merge, production gated behind explicit approval, branch protection rules — driving **30% reduction in deployment cycle time** and **zero broken-build incidents** over four months.
- Delivered the **Automated Content Intelligence Pipeline**: Puppeteer-based headless scraping (`waitUntil: 'networkidle2'` for SPA hydration), TypeScript normalization, dual-model AI rewriting (Gemini primary →

OpenAI failover per article), React 19 Diff View for human-in-the-loop editorial approval, Dockerized microservices on GitHub Actions cron schedule.

- Shipped a parallel **Website Builder product** within a 6-week window at zero headcount increase by auditing and reusing **~60% of shared infrastructure** (auth, Supabase layer, Razorpay billing, CI/CD); restructured sprint cadence to separate greenfield work from maintenance load.
- Implemented **idempotent Razorpay subscription billing**: HMAC webhook validation, payment-state reconciliation, out-of-order event handling via re-entrant idempotency check — achieving **100% payment-state consistency** across all paying customers.

Vellore Institute of Technology | Research Contributor | Jan 2025 – Nov 2025 | Vellore (On-campus)

- Contributed to research on **Blockchain-Based LLM Model Using Fully Homomorphic Encryption (FHE) for Academic Records** — investigating how FHE enables computation on encrypted student data without decryption, with blockchain providing an immutable audit trail for credential issuance and verification.
- Evaluated architectural trade-offs between on-chain data storage and off-chain encrypted record pointers, assessing gas cost constraints versus data privacy guarantees in an academic credential provenance system.
- Explored how model inference could operate on privacy-preserving representations of sensitive academic data using FHE-compatible cryptographic primitives, without exposing plaintext records.

PROJECTS

Nimbus — AI-First RAG Document Workspace (in progress)

Stack: Next.js 14 · TypeScript · PostgreSQL · pgvector · Prisma · NextAuth.js · Groq API · OpenAI

- Production-grade RAG pipeline: async queue-based ingestion (status states: pending → processing → indexed → error), overlap chunking to prevent semantic loss at chunk boundaries, embedding generation, and storage in a pgvector extension column within PostgreSQL.
- **Cosine similarity search** via pgvector (embedding $\leftrightarrow$ query_vector) co-locates vector data with relational metadata (user_id, document_id) for hybrid SQL + semantic filtering — **sub-800ms retrieval** target.
- **Provider-agnostic LLM abstraction layer**: single config flag (provider: "groq" | "openai" | "anthropic") routes generation calls — switching inference providers requires zero code changes.
- **Streaming responses** via Groq API delivered as SSE, providing real-time token-by-token document Q&A. Ingestion pipeline includes **429 retry logic** (exponential backoff) and per-document status tracking.

Autonomous Web Agent — Claude 3.5 Sonnet + MCP Orchestration

GitHub: github.com/atinsharma24/auto-agent · Stack: Python · Groq API · BrowserOS (MCP) · OpenClaw · Claude 3.5 Sonnet

- **Dual-agent architecture**: Executor (Claude 3.5 Sonnet) handles complex multi-step DOM + form reasoning; Watcher (Groq fast inference) validates page state between every transition — separating expensive reasoning from high-frequency cheap polling.
- **MCP (Model Context Protocol)** via BrowserOS gives the Executor structured tool-level access to browser DOM (inspect labels, types, values), form-fill, and navigation — intent-driven traversal with no brittle CSS selectors.
- Persistent **local queue manager** (status: pending → in_progress → applied / failed / escalated) with URL deduplication; **WAF/CAPTCHA detection** via DOM fingerprint — Watcher escalates to human, Executor resumes from last checkpoint.

DocGPT — RAG Document Assistant

Stack: React 19 · Python (Django) · ChromaDB · OpenAI GPT-4o-mini · all-MiniLM-L6-v2 · PyPDF2 · Docker

- Full Django RAG pipeline: PDF extraction via PyPDF2 with overlap chunking (~500–1000 tokens, ~50–100 overlap), **local embedding** via all-MiniLM-L6-v2 (384-dimensional, offline — zero API cost), ChromaDB cosine retrieval, GPT-4o-mini generation.
- Containerized Django backend + ChromaDB as Docker services; prompt structure: system context + k-nearest retrieved chunks + user query for grounded generation.

VyaparGPT — WhatsApp AI Business Assistant (OpenBiz product)

Stack: Node.js · OpenAI API · Google Gemini API · WhatsApp Business API

- Full WhatsApp message lifecycle: webhook ingestion → HMAC-SHA256 validation → message-type routing (text / media / document / status) → session context → LLM dispatch → delivery.
- Session-based multi-turn context (rolling conversation_history[]) with configurable TTL expiry; Gemini Vision API extracts structured fields from photographed SMB documents (GST certificates, invoices, vendor agreements).

Automated Content Intelligence Pipeline (OpenBiz product)

Stack: React 19 · TypeScript · Laravel · Node.js · OpenAI · Google Gemini · Docker · GitHub Actions · Puppeteer

- Puppeteer headless Chromium (waitUntil: 'networkidle2') for SPA scraping; dual-model AI rewriting (Gemini primary → OpenAI fallback per article); React 19 split-screen Diff View with per-article state machine (pending_review → approved → published).
- Dockerized microservices (Scraper, AI Rewriter, Laravel API, React Frontend) with unless-stopped restart policy; GitHub Actions cron schedule drives fully automated ingestion and review cycles.

TECHNICAL SKILLS

Languages	JavaScript, TypeScript, Python, SQL, Java
Frontend	React.js (v18, v19), Next.js, React Native, Redux Toolkit, TailwindCSS
Backend	Node.js, Express.js, NestJS, Django, Laravel
Databases	PostgreSQL, Supabase, MongoDB, ChromaDB, pgvector, AWS DynamoDB
AI / LLM	OpenAI API, Google Gemini API, Groq API, RAG Pipelines, pgvector, ChromaDB, Sentence Transformers (all-MiniLM-L6-v2), Vercel AI SDK, LangChain.js, Prompt Engineering, PyPDF2, Claude 3.5 Sonnet (Anthropic API), MCP
DevOps / Cloud	Docker, GitHub Actions, Vercel, AWS S3, AWS SNS
Tools	Git, Postman, Socket.IO, Razorpay SDK, VideoSDK, Puppeteer, BrowserOS / OpenClaw
Patterns	WebSocket (Socket.IO), REST, Webhook Validation (HMAC-SHA256), N+1 Query Optimization, Vector Similarity Search (cosine), Provider-Agnostic LLM Layers, Idempotent Event Processing

EDUCATION

Vellore Institute of Technology, Vellore

B.Tech in Computer Science and Engineering | Graduating May 2026

Focus: Full Stack Development, AI Systems

Delhi Public School, Agra

Secondary & Senior Secondary

PROFILE SNAPSHOT

Availability	Immediate Joiner — 0-day notice period
Work Mode	Remote-first; open to Bangalore, Hyderabad, Pune, Delhi NCR, Mumbai
Expected CTC	15–18 LPA (INR) — negotiable for high-ownership roles with equity
Work Authorization	Indian Citizen — no sponsorship required
GitHub	github.com/atinsharma24/
LinkedIn	linkedin.com/in/atinsharma24/

