

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Full-Stack Product Engineer with founding-team experience building AI systems for Indian SMBs.

- Shipped VyaparGPT — WhatsApp-native LLM assistant (40 SMB pilot)
- Reduced API latency by 40% via PostgreSQL optimization (N+1 fixes, caching)
- Built production RAG pipelines (pgvector, sub-800ms retrieval)

PROFESSIONAL EXPERIENCE

OpenBiz Software India Pvt Ltd

Founding Engineer

Jul 2025 – Mar 2026

Remote/Bangalore

- Architected the zero-to-one core MERN platform, implementing strict TypeScript practices across the stack to build a scalable, production-ready foundation.
- Architected VyaparGPT, a WhatsApp-native AI assistant for Indian SMBs, successfully deploying the MVP to an initial closed-pilot cohort of ~40 businesses; handled webhook ingestion → LLM orchestration → response delivery with provider fallback and session-aware context.
- Built an AI document verification module using Google Gemini API, automating data extraction from unstructured SMB documents and replacing a manual review step.
- Diagnosed and resolved N+1 query patterns in Supabase (PostgreSQL), reducing API latency by 40% across critical endpoints through indexing and response caching.
- Engineered a fault-tolerant Razorpay subscription billing system with idempotent webhook validation, achieving 100% payment-state consistency and reducing manual reconciliation to zero.
- Owned CI/CD pipeline setup using GitHub Actions and Vercel, reducing deployment cycle time by 30% and eliminating environment-drift bugs.

PROJECTS

Nimbus — AI RAG Workspace | github.com/atinsharma24/nimbus

In Development 2026

- Stack: Next.js, pgvector, Groq, Prisma
- Architected a production RAG pipeline (pgvector + Groq API) achieving sub-800ms semantic retrieval over multi-format document corpora with streaming LLM responses.
- Designed a provider-agnostic LLM layer supporting Groq, OpenAI, and Anthropic, enabling one-config switching without downstream changes to retrieval logic.
- Implemented async ingestion with 429 retry logic and per-document status tracking for production-grade rate limit handling.

Autonomous Web Agent | github.com/atinsharma24/auto-agent

- Stack: OpenClaw, BrowserOS (MCP), Python, Groq
- Designed a Claude 3.5 Sonnet + MCP orchestration layer that translates structured local context schemas into DOM-level form interactions for multi-step web workflows across dynamic web forms.
- Built a decoupled system with a lightweight Groq-based "Watcher" for cost-efficient parsing and a heavy "Executor" agent for complex interactions.
- Engineered resilient infra with local queue management and WAF/Captcha detection with graceful human-escalation.

Automated Content Intelligence Pipeline | github.com/atinsharma24/pipeline

- Stack: React 19, TypeScript, Node.js, OpenAI
- Built an autonomous content pipeline running Puppeteer scraping and dual-model AI rewriting with Gemini and OpenAI failover.
- Developed a React 19 frontend with a split-screen diff view for real-time editorial comparison of AI-rewritten content.

SKILLS

Languages: JavaScript/TypeScript, Python, SQL, Java

Frontend: React (v18/19), Next.js, React Native, Redux Toolkit, TailwindCSS

Backend/DB: Node.js, Express, NestJS, Django, PostgreSQL, MongoDB, pgvector, ChromaDB

AI/LLM: OpenAI, Google Gemini, Groq API, RAG Pipelines, LangChain.js, Sentence Transformers, Prompt Engineering

DevOps: Docker, GitHub Actions, Vercel, AWS (S3, DynamoDB)

Tools: Git, Postman, Socket.IO, Razorpay SDK, VideoSDK, Puppeteer

EDUCATION

Vellore Institute of Technology

B.Tech in Computer Science and Engineering

- Focus on Full Stack Development and AI Systems.

Expected Graduation: May 2026

Vellore