

ATIN SHARMA

Agra, India | atinsharma24@gmail.com | +91 82185 02886
linkedin.com/in/atinsharma24 | github.com/atinsharma24

PROFILE

Applied AI Engineer with production experience building agentic systems, RAG pipelines, and MCP server infrastructure. Shipped real-user AI products (LLM orchestration, WhatsApp automation, document extraction) and built agent runtime systems handling state management, retries, and structured LLM outputs. Strong Python backend foundation with hands-on production MCP orchestration, pgvector semantic search, and provider-agnostic LLM layers across OpenAI, Groq, and Gemini.

- Built production MCP orchestration layer (Claude 3.5 + BrowserOS MCP) for multi-step agentic workflows with state management and human-escalation.
- Architected Nimbus AI: provider-agnostic RAG workspace with async ingestion, pgvector cosine retrieval, and streaming Groq responses.
- Shipped VyaparGPT (WhatsApp AI assistant): full LLM orchestration pipeline with session-aware context and provider-fallback to 1,000+ SMB users.

PROFESSIONAL EXPERIENCE

OpenBiz Software India Pvt Ltd

Jul 2025 – Mar 2026

Founding Engineer

Remote/Bangalore

- Architected VyaparGPT, a WhatsApp-native AI assistant: designed LLM orchestration pipeline from webhook ingestion → session context management → structured output generation → response delivery; implemented OpenAI-to-Gemini provider failover for production reliability.
- Built an AI document verification agent using Google Gemini API, extracting structured data from unstructured SMB documents (GST certificates, vendor agreements) and eliminating a manual review step.
- Implemented OAuth integrations and SaaS webhook systems (WhatsApp Business API, Razorpay) with idempotent validation, achieving 100% payment-state consistency.
- Diagnosed and resolved N+1 query patterns in Supabase (PostgreSQL), reducing API latency by 40% through composite indexing and TTL-based response caching.
- Owned CI/CD pipeline (GitHub Actions + Vercel), cutting deployment cycle time by 30% and eliminating environment-drift incidents.

PROJECTS

Autonomous Web Agent | Production MCP Server

2026

- Stack: Python, Claude 3.5 Sonnet, BrowserOS MCP, Groq API
- Designed and deployed a production MCP server + Claude orchestration layer for multi-step agentic web workflows; translates structured schemas into DOM-level form interactions across dynamic portals.
- Implemented agent runtime with local queue management, retry logic, state persistence, and WAF/Captcha detection with graceful human-escalation fallback.
- Decoupled architecture: lightweight Groq-based "Watcher" agent for cost-efficient parsing; heavy "Executor" agent for complex multi-turn interactions.

Nimbus – AI RAG Workspace | github.com/atinsharma24/nimbus

In Development 2026

- Stack: Next.js, Python, pgvector, Groq, Prisma, PostgreSQL
- Built production RAG pipeline (pgvector + Groq) with sub-800ms semantic retrieval, streaming LLM responses, and per-document ingestion status tracking.
- Designed provider-agnostic LLM layer (Groq, OpenAI, Anthropic) with one-config switching; built evaluation harness for retrieval accuracy across document corpora.
- Async ingestion with 429 retry logic, exponential backoff, and structured output schemas for production-grade rate-limit handling.

Automated Content Intelligence Pipeline

- Stack: Node.js, Python, OpenAI, Gemini, Docker, GitHub Actions, Puppeteer
- Built an autonomous agent pipeline with Puppeteer scraping, dual-model AI rewriting (Gemini + OpenAI failover), and microservices on Docker with scheduled GitHub Actions cron execution.

SKILLS

Agent / AI: MCP Servers, LLM Orchestration, RAG Pipelines, Agentic Workflows, Prompt Engineering, Structured Outputs, Eval Frameworks, pgvector, ChromaDB

Languages: Python, JavaScript/TypeScript, SQL

Backend/DB: Node.js, Express, NestJS, Django, PostgreSQL, Redis, MongoDB, pgvector

LLM APIs: OpenAI, Google Gemini, Groq API, Anthropic Claude, LangChain.js, Vercel AI SDK

Integrations: OAuth 2.0, WhatsApp Business API, Razorpay, REST Webhooks, SaaS API connectors

DevOps: Docker, GitHub Actions, Vercel, AWS S3/DynamoDB

EDUCATION

Vellore Institute of Technology

B.Tech in Computer Science and Engineering

Coursework Completed: Nov 2025 | Degree Conferred: May 2026

Vellore